

Financial Mathematics 32000

Lecture 5

Roger Lee

2025 April 21

UNIT 3: Monte Carlo

Generating random variables

Monte Carlo estimate

Let Y be a discounted payoff. Example: $Y = e^{-\int_0^T r_t dt} (S_T - K)^+$.

Want to calculate the time-0 price $C = \mathbb{E}Y$

Generate Y_1, Y_2, \dots independently and identically distributed as Y .

(How?)

Then the random variable

$$\hat{C}_M := \frac{Y_1 + Y_2 + \dots + Y_M}{M}$$

is the Monte Carlo estimate of C . Note that

$$\mathbb{E}\hat{C}_M = C$$

is the expectation of the price,
equal to the true price?

By the strong law of large numbers, with probability 1 we have

$$\hat{C}_M \rightarrow C \text{ as } M \rightarrow \infty.$$

as we do more simulations, the
simulated discounted payoff will
converge to the price

we use MC because we don't know the expected payout and the var of the payout

How fast does convergence occur?

Let $\sigma^2 := \text{Var}(Y)$. (Here σ does not denote volatility.) Then

$$\text{Var}(\hat{C}_M) = \frac{1}{M^2} \text{Var}(Y_1 + Y_2 + \cdots + Y_M) = \frac{1}{M^2} (M\sigma^2) = \frac{\sigma^2}{M}$$

By the Central Limit Theorem, we have convergence in distribution

$$\frac{\hat{C}_M - \mathbb{E}\hat{C}_M}{\sqrt{\text{Var } \hat{C}_M}} \xrightarrow{d} N(0, 1) \quad \text{hence} \quad \frac{\hat{C}_M - C}{\sigma/\sqrt{M}} \xrightarrow{d} N(0, 1)$$

as $M \rightarrow \infty$. Conclusion still holds using sample stdev in place of σ :

$$\frac{\hat{C}_M - C}{\hat{\sigma}_M/\sqrt{M}} \xrightarrow{d} N(0, 1)$$

because $\hat{\sigma}_M^2 \rightarrow \sigma^2$ where

$$\hat{\sigma}_M^2 := \frac{1}{M-1} \sum_{m=1}^M (Y_m - \hat{C}_M)^2.$$

is the “sample variance” and $\hat{\sigma}_M$ is the “sample standard deviation.”

Confidence intervals

For 95% confidence interval, $p=0.05$, $C_M + 1.96 \cdot \text{std_error}$

So an asymptotic (for large M) confidence interval of $100(1-p)\%$ is

$$\left(\hat{C}_M - \underbrace{\mathcal{N}^{-1}\left(1 - \frac{p}{2}\right)}_{1.96} \frac{\hat{\sigma}_M}{\sqrt{M}}, \hat{C}_M + \underbrace{\mathcal{N}^{-1}\left(1 - \frac{p}{2}\right)}_{1.96} \frac{\hat{\sigma}_M}{\sqrt{M}} \right)$$

where \mathcal{N} is the standard Normal cdf.

- ▶ $\mathcal{N}^{-1}\left(1 - \frac{p}{2}\right)$ tells us: a radius of *how many standard deviations* of a normal distribution contains $100(1-p)\%$ of the probability? **1.96**
- ▶ The $\hat{\sigma}_M/\sqrt{M}$ is called the **standard error**. **how many \$ is one stdev**
It gives us the estimated standard deviation of \hat{C}_M .

- ▶ Example: Let $p = 0.05$. Then $\mathcal{N}^{-1}\left(1 - \frac{p}{2}\right) \approx 1.96$ and

$$\left(\hat{C}_M - 1.96 \frac{\hat{\sigma}_M}{\sqrt{M}}, \hat{C}_M + 1.96 \frac{\hat{\sigma}_M}{\sqrt{M}} \right)$$

is an asymptotic 95% confidence interval for C .

Confidence intervals

- ▶ If $\hat{\sigma} = 20$ and we run $M = 10000$ simulations, then a 95% confidence interval has radius

$$1.96 \times \frac{20}{\sqrt{10000}} = 0.40 \quad \text{radius of 40 cents}$$

To reduce this to 0.04, we need to take $M = 1$ million.

radius of 4 cents

- ▶ So we will want to use *variance reduction* techniques, which reformulate the problem to keep σ small, or which carefully introduce dependence in the simulations to keep $\text{Var } \hat{C}_M$ small.

UNIT 3: Monte Carlo

Generating random variables

If we want the 23rd percentile, take the F^{-1} of 0.23 and we will get a number X that is smaller than 0

If we want the 88th percentile, take the F^{-1} of 0.88 and we will get a number X that is bigger than 0

Look at notebook

The inverse CDF method

- ▶ Assume the existence of a pseudo-random number generator whose output can be treated as if it is IID uniform on $(0, 1)$.
Python: `numpy.random.Generator` has method `random()`
- ▶ To generate a random variable X having a CDF F , generate $U \sim \text{Uniform}(0, 1)$, then apply F^{-1} , the inverse CDF, to produce

$$X := F^{-1}(U)$$

As desired, $\mathbb{P}(X \leq x) = \mathbb{P}(F^{-1}(U) \leq x) = \mathbb{P}(U \leq F(x)) = F(x)$.

(If F not invertible, letting $F^{-1}(u) := \min\{x : F(x) \geq u\}$ works.)

Intuition: randomly choose the *percentile* between 0% and 100%, uniformly. The inverse CDF finds the corresponding value of X .

Generating normal random variables

- ▶ Python: `numpy.random.Generator` has method `normal()`

But what if you need to build your own?

- ▶ Could do $\mathcal{N}^{-1}(U)$, if an implementation of the inverse of the normal CDF \mathcal{N} is available. Excel: `NORMSINV(RAND())`
- ▶ Box-Muller method: If (X, Y) are independent $\text{Normal}(0, 1)$, then $R := X^2 + Y^2$ has CDF $\mathbb{P}(R \leq r) = 1 - e^{-r/2}$. Given R , the point (X, Y) is uniformly distributed on the circle of radius \sqrt{R} . So generate *pairs* of independent normals by drawing U_1 and U_2 IID from a $\text{Uniform}(0, 1)$ distribution, and taking

SKIP BOX MULLER

$$R := -2 \log(U_1)$$

$$(X, Y) := (\sqrt{R} \cos(2\pi U_2), \sqrt{R} \sin(2\pi U_2))$$

Simulating random variables in Python

default Random Number Generator

```
[1] import numpy
    rng = numpy.random.default_rng(seed=0)
```

By starting w/ the same seed everytime, we will generate the same random numbers

```
[2] rng.random(size=5)
```

```
array([0.63696169, 0.26978671, 0.04097352, 0.01652764, 0.81327024])
```

```
[3] rng.random(size=5)    if we put only random, it will be between 0 and 1
```

```
array([0.91275558, 0.60663578, 0.72949656, 0.54362499, 0.93507242])
```

```
[4] rng.normal(size=5)    Normal dist, no bounded domain
```

```
array([-0.62327446,  0.04132598, -2.32503077, -0.21879166, -1.24591095])
```

```
[5] rng.normal(size=(2,5))
```

```
array([[ -0.73226735, -0.54425898, -0.31630016,  0.41163054,  1.04251337],
       [-0.12853466,  1.36646347, -0.66519467,  0.35151007,  0.90347018]])
```

In some cases, no need to simulate entire path

If Y is a known function of random variables with distributions that you can directly simulate, then it is easy to generate Y_1, Y_2, \dots

Example: Consider a call paying $(S_T - K)^+$ where S is GBM.

$$dS_t = rS_t dt + \sigma S_t dW_t$$

Then

$$Y = e^{-rT}(S_T - K)^+ = e^{-rT}(S_0 e^{(r-\sigma^2/2)T + \sigma W_T} - K)^+$$

where $W_T \sim N(0, T)$. So let

the m -th simulation being run

$$Y_m := e^{-rT}(S_T - K)^+ = e^{-rT}(S_0 e^{(r-\sigma^2/2)T + \sigma\sqrt{T}Z^{(m)}} - K)^+$$

where the $Z^{(m)}$ are IID standard normal: $N(0, 1)$.

But sometimes need to simulate entire path

- ▶ But what about more complicated dynamics, such as

$$dS_t = rS_t dt + \sigma(S_t, t)S_t dW_t$$

here, sigma is not constant, but depends on the underlying

- ▶ Or what if σ or r follows a process driven by a second BM.
- ▶ What about more complicated contracts, such as an Asian option or a barrier option.


We may need to simulate the whole path.

Geometric Brownian motion: $dX_t = \mu X_t dt + \sigma X_t dW_t$


notebook: $\text{dlog}(X_t) = \dots$

Let $X_0 = 100$. Trajectories for $\mu = -0.15, +0.15$ and $\sigma = 0.20, 0.40$:


miu = 0.15
sigma = 0.2




miu = 0.15
sigma = 0.4



miu = -0.15
sigma = 0.2



miu = -0.15
sigma = 0.4



Euler method to simulate path of a state variable

Suppose X satisfies

In one segment (ex. region I), we are fixing a and b
 But the moment we go to the next time step, a new segment, region II,
 we have a new a and b
 In the same region, between t_n and t_{n+1} , we have a fixed a and b
 Look at notebook

$$dX_t = a(X_t, t)dt + b(X_t, t)dW_t$$

Divide the time interval $[0, T]$ into N parts: $\Delta t = T/N$, $t_n = n\Delta t$.

Define the m th simulated path by initializing $X_0^{(m)} = X_0$,

and given X_{t_n} , obtain $X_{t_{n+1}}$ by

$$X_{t_{n+1}}^{(m)} = X_{t_n}^{(m)} + a(X_{t_n}^{(m)}, t_n)\Delta t + b(X_{t_n}^{(m)}, t_n)Z_n^{(m)}\sqrt{\Delta t} \quad dW_t$$

where $Z_n^{(m)}$ are IID standard normal. Evaluate the discounted payoff for the m th path. Take the average across all paths $m = 1, \dots, M$.

This extends directly to multidimensional state vectors X and multidimensional standard Brownian motion W .

Euler method convergence

With some assumptions on a and b , the Euler method has *weak* order of convergence 1 for general f , meaning

SKIP

$$|\mathbb{E}f(X_T) - \mathbb{E}f(X_T^{(m)})| = O(\Delta t)$$

(Estimating $\mathbb{E}f(X_T^{(m)})$ produces additional error, not included here.)

Error analysis: Let $a_t = a(X_t)$ and $b_t = b(X_t)$. First subinterval:

$$\begin{aligned} X_{t_1} &= X_0 + \int_0^{t_1} a_t dt + \int_0^{t_1} b_t dW_t \\ &= X_0 + \int_0^{t_1} \left(a_0 + \int_0^t a \frac{\partial a}{\partial x} + \frac{1}{2} b^2 \frac{\partial^2 a}{\partial x^2} ds + \int_0^t b \frac{\partial a}{\partial x} dW_s \right) dt \\ &\quad + \int_0^{t_1} \left(b_0 + \int_0^t a \frac{\partial b}{\partial x} + \frac{1}{2} b^2 \frac{\partial^2 b}{\partial x^2} ds + \int_0^t b \frac{\partial b}{\partial x} dW_s \right) dW_t \end{aligned}$$

The Euler scheme keeps **three terms** to generate $X_{t_1}^{(m)}$

Crude analysis of discretization error of Euler scheme

SKIP

- ▶ Intuition for weak error: $|\mathbb{E}(X_T - X_T^{(m)})|$ is $O(\Delta t)$, because there are N time steps, and \mathbb{E} of error is $O(\Delta t)^2$ at each step (the only nonzero- \mathbb{E} error term in previous equation is the $dsdt$ term).
- ▶ Euler scheme has strong order of convergence $1/2$, meaning that

$$\mathbb{E}|X_T - X_T^{(m)}| = O(\Delta t)^{1/2} \quad \text{as } \Delta t \rightarrow 0$$

Intuition for strong error: Biggest ignored term is $dW_s dW_t$.

Variance of error at one time step = $O(\text{Var}(\Delta W)^2) = O(\Delta t)^2$.

Variance of total error $X_T - X_T^{(m)}$ after N time steps = $O(\Delta t)$.

Standard deviation of $X_T - X_T^{(m)}$ after N time steps = $O(\Delta t)^{1/2}$.

This suggests that strong order of convergence is $1/2$.

Milstein scheme

SKIP

Milstein scheme: Don't ignore the term

$$\int_0^{t_1} \int_0^t b(X_s) \frac{\partial b}{\partial x}(X_s) dW_s dW_t.$$

Approximate it as

$$\begin{aligned} b(X_0) \frac{\partial b}{\partial x}(X_0) \int_0^{t_1} \int_0^t dW_s dW_t &= b(X_0) \frac{\partial b}{\partial x}(X_0) \int_0^{t_1} W_t dW_t \\ &= \frac{1}{2} b(X_0) \frac{\partial b}{\partial x}(X_0) (W_{t_1}^2 - t_1) \end{aligned}$$

So Milstein is

$$X_{t_{n+1}}^{(m)} = X_{t_n}^{(m)} + a(X_{t_n}^{(m)}) \Delta t + b(X_{t_n}^{(m)}) Z_n^{(m)} \sqrt{\Delta t} + \frac{1}{2} b \frac{\partial b}{\partial x}(X_{t_n}^{(m)}) ([Z_n^{(m)}]^2 - 1) \Delta t$$

Milstein has strong order of convergence 1, and weak order 1.

Weak convergence (important in option pricing): same order as Euler.

Covariance and correlation matrices

Recall that the covariance matrix of a zero-mean vector Z is $\mathbb{E}(ZZ^\top)$.

Let M be a real symmetric matrix. The following are equivalent:

- ▶ M is a covariance matrix of some vector.
- ▶ M is *positive semi-definite*, which means that $x^\top Mx \geq 0$ for all real vectors x .
- ▶ The eigenvalues of M are all nonnegative.
- ▶ The principal minors of M are all nonnegative. (Principal minors = the determinants of the matrices formed by crossing out any rows and corresponding columns of M).
- ▶ M has a *Cholesky decomposition* $LL^\top = M$ for some real lower-triangular matrix L with diagonal entries ≥ 0 .

Covariance and correlation matrices

Moreover, the following are equivalent:

- ▶ M is a correlation matrix
- ▶ M is a covariance matrix and its diagonal elements are all 1.

Moreover, if M is a 3×3 matrix, the following are equivalent

- ▶ M is a correlation matrix.
- ▶ M is symmetric, its entries $\in [-1, 1]$, its diagonal entries $= 1$, and $\det M \geq 0$.

$$\text{Corr}(x,y) = \text{Cov}(x/\text{stdev}(x), y/\text{stdev}(y))$$

Interview question

Suppose $\text{Corr}(X, Y) = \text{Corr}(X, Z) = \text{Corr}(Y, Z) = \rho$.

What are the possible values of ρ ?

Notebook

Generating correlated Brownian motions

To get a D -dimensional vector \bar{W} of BM with correlation matrix H , find a matrix $L \in \mathbb{R}^{D \times D}$ such that $LL^\top = H$, and let $\bar{W} = LW$, where W is standard BM in \mathbb{R}^D .

So the process $dX_t = a(X_t, t)dt + b(X_t, t)d\bar{W}_t$ becomes

$$dX_t = a(X_t, t)dt + b(X_t, t)L dW_t$$

Simulation by Euler method:

$$X_{t_{n+1}}^{(m)} = X_{t_n}^{(m)} + a(X_{t_n}^{(m)}, t_n)\Delta t + b(X_{t_n}^{(m)}, t_n)L Z_n^{(m)}\sqrt{\Delta t}$$

where $Z_n^{(m)}$ are IID standard normal in \mathbb{R}^D . How to find L ?

- ▶ `numpy.random.Generator.multivariate_normal` generates LZ
- ▶ `numpy.linalg.cholesky` returns L given H

Generating correlated Brownian motions

Example: if we want $\bar{W} = \begin{pmatrix} W^{[1]} \\ W^{[2]} \end{pmatrix}$ with $\text{corr}(\Delta W^{[1]}, \Delta W^{[2]}) = \rho$, then

$$H = \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix}$$

and Cholesky finds $LL^\top = H$ where

$$L = \begin{pmatrix} 1 & 0 \\ \rho & \sqrt{1 - \rho^2} \end{pmatrix}$$

If Cholesky routine unavailable, can solve for L by traversing [the upper or lower triangular part of] H entry-by-entry. Each entry gives rise to an equation involving elements of L and only one unknown.

Generating correlated Brownian motions

Sometimes it is not necessary to simulate the entire path.

- ▶ Suppose $X_T - X_0$ is known to be multivariate normal with mean μT and covariance matrix HT .
- ▶ Suppose the option payoff depends only on X_T .

Then no need to divide $[0, T]$ into N steps. No need for Euler.

- ▶ Just generate

$$X_T^{(m)} := X_0 + \underbrace{\mu T}_{Dx1} + \underbrace{L}_{DxD} \underbrace{\sqrt{T}}_{Dx1} Z^{(m)}$$

where $LL^\top = H$ and the $Z^{(m)}$ are IID standard normal in \mathbb{R}^D .